


A General Recipe for Likelihood-free Bayesian Optimization

Jiaming Song*, Lantao Yu*, Willie Neiswanger, Stefano Ermon



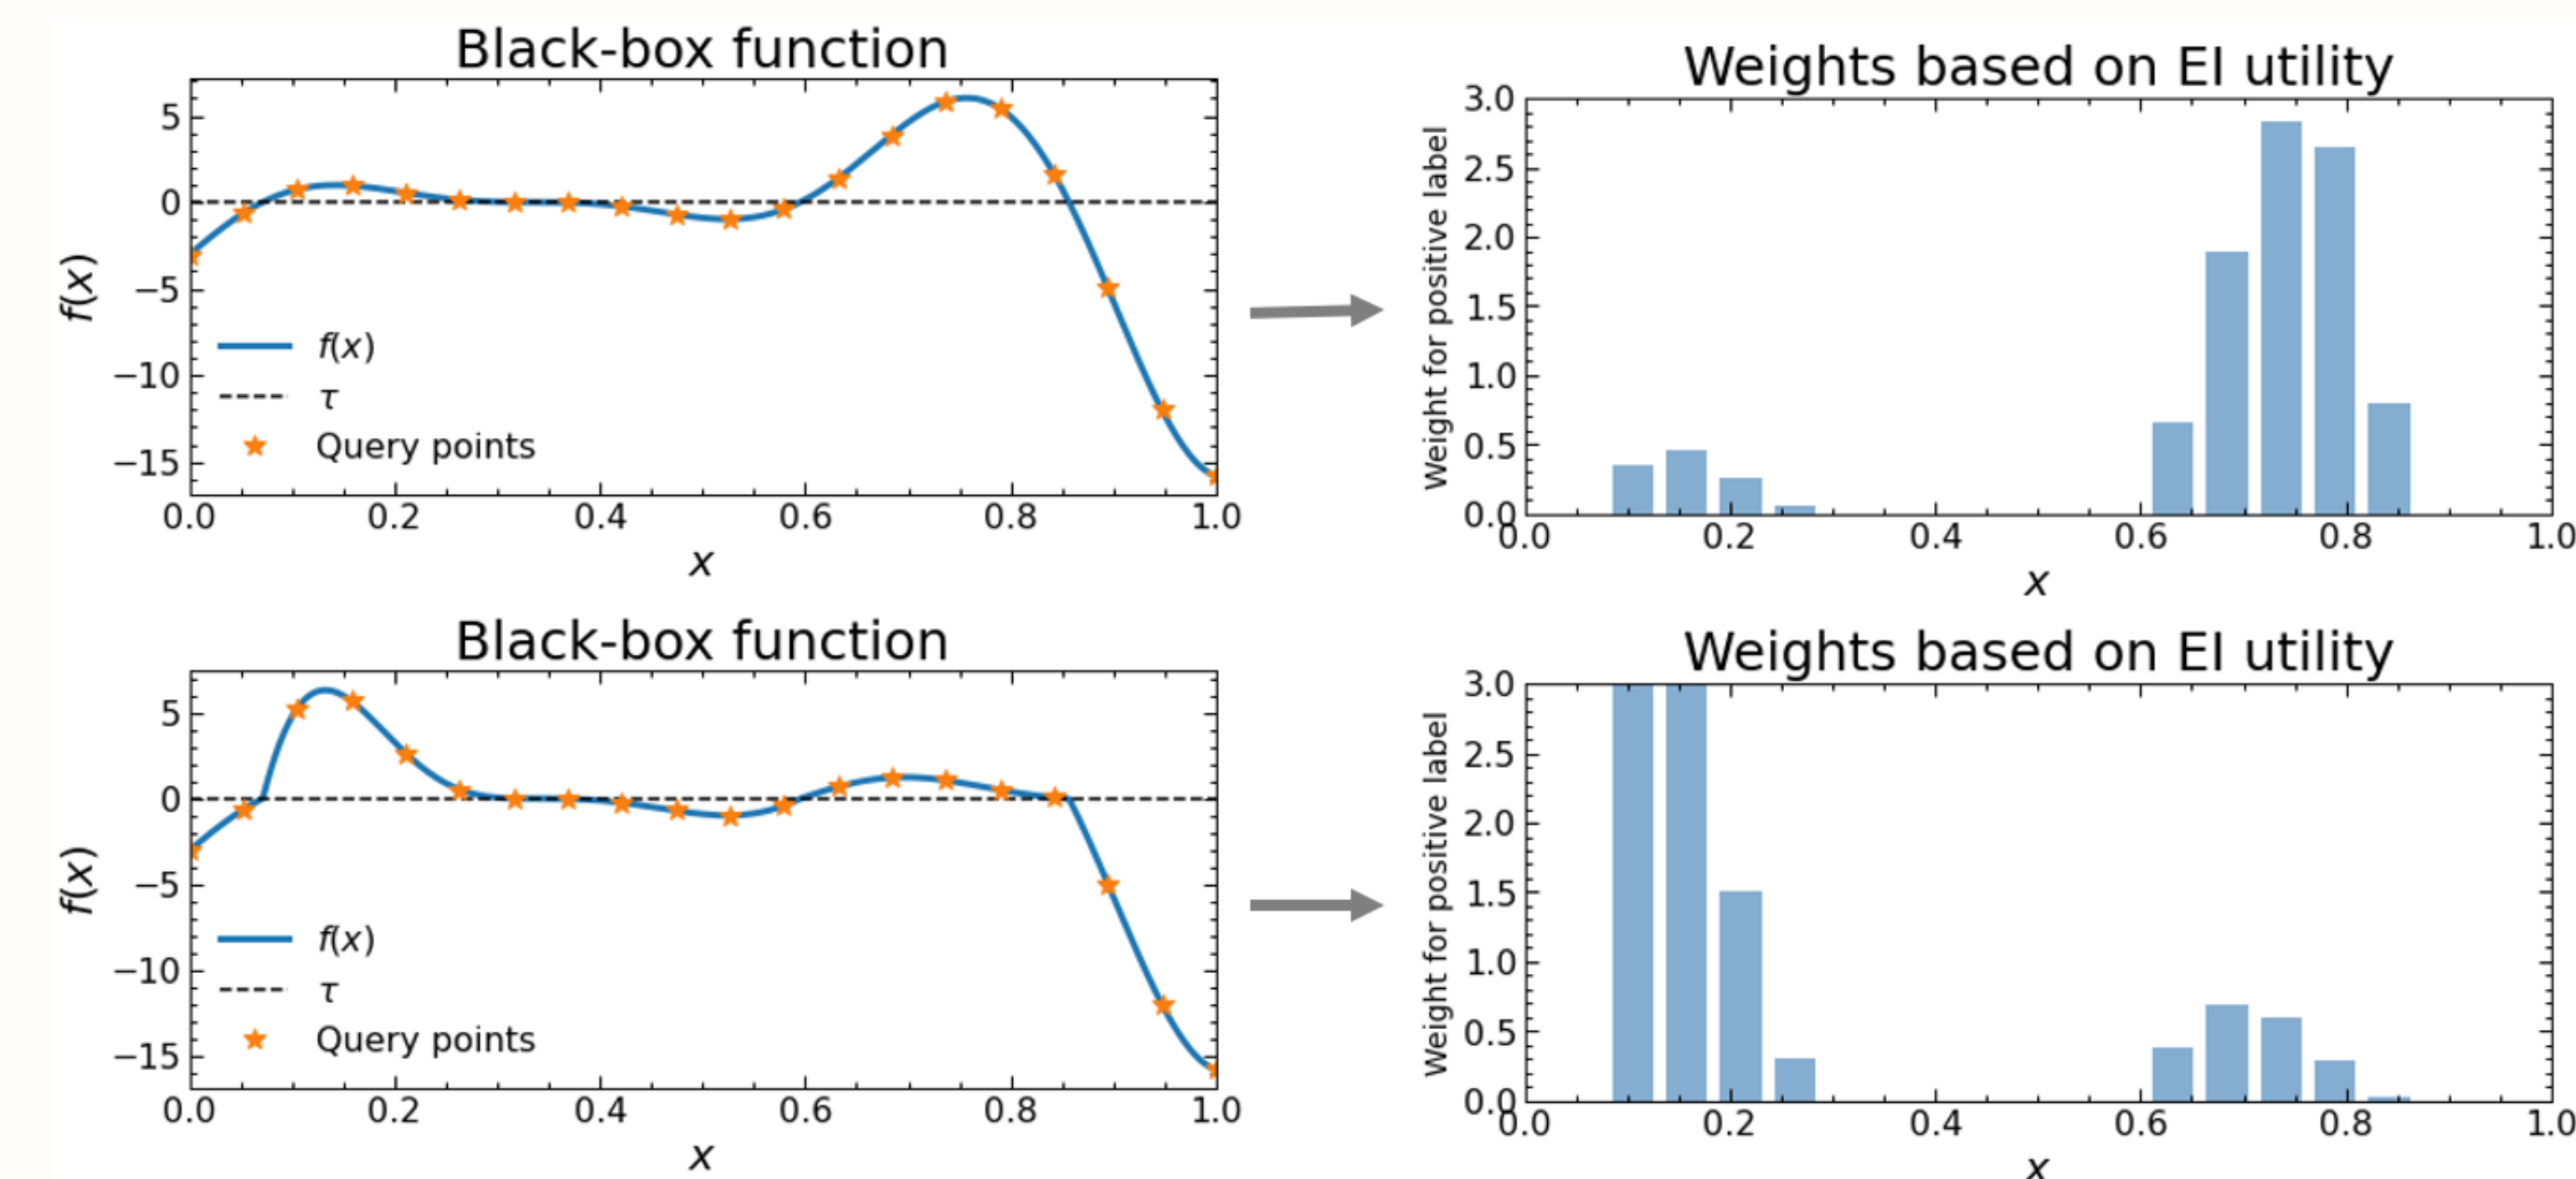
BO for Global Optimization of Black-box Functions

Function evaluations are expensive, e.g. neural architecture search and material discovery.



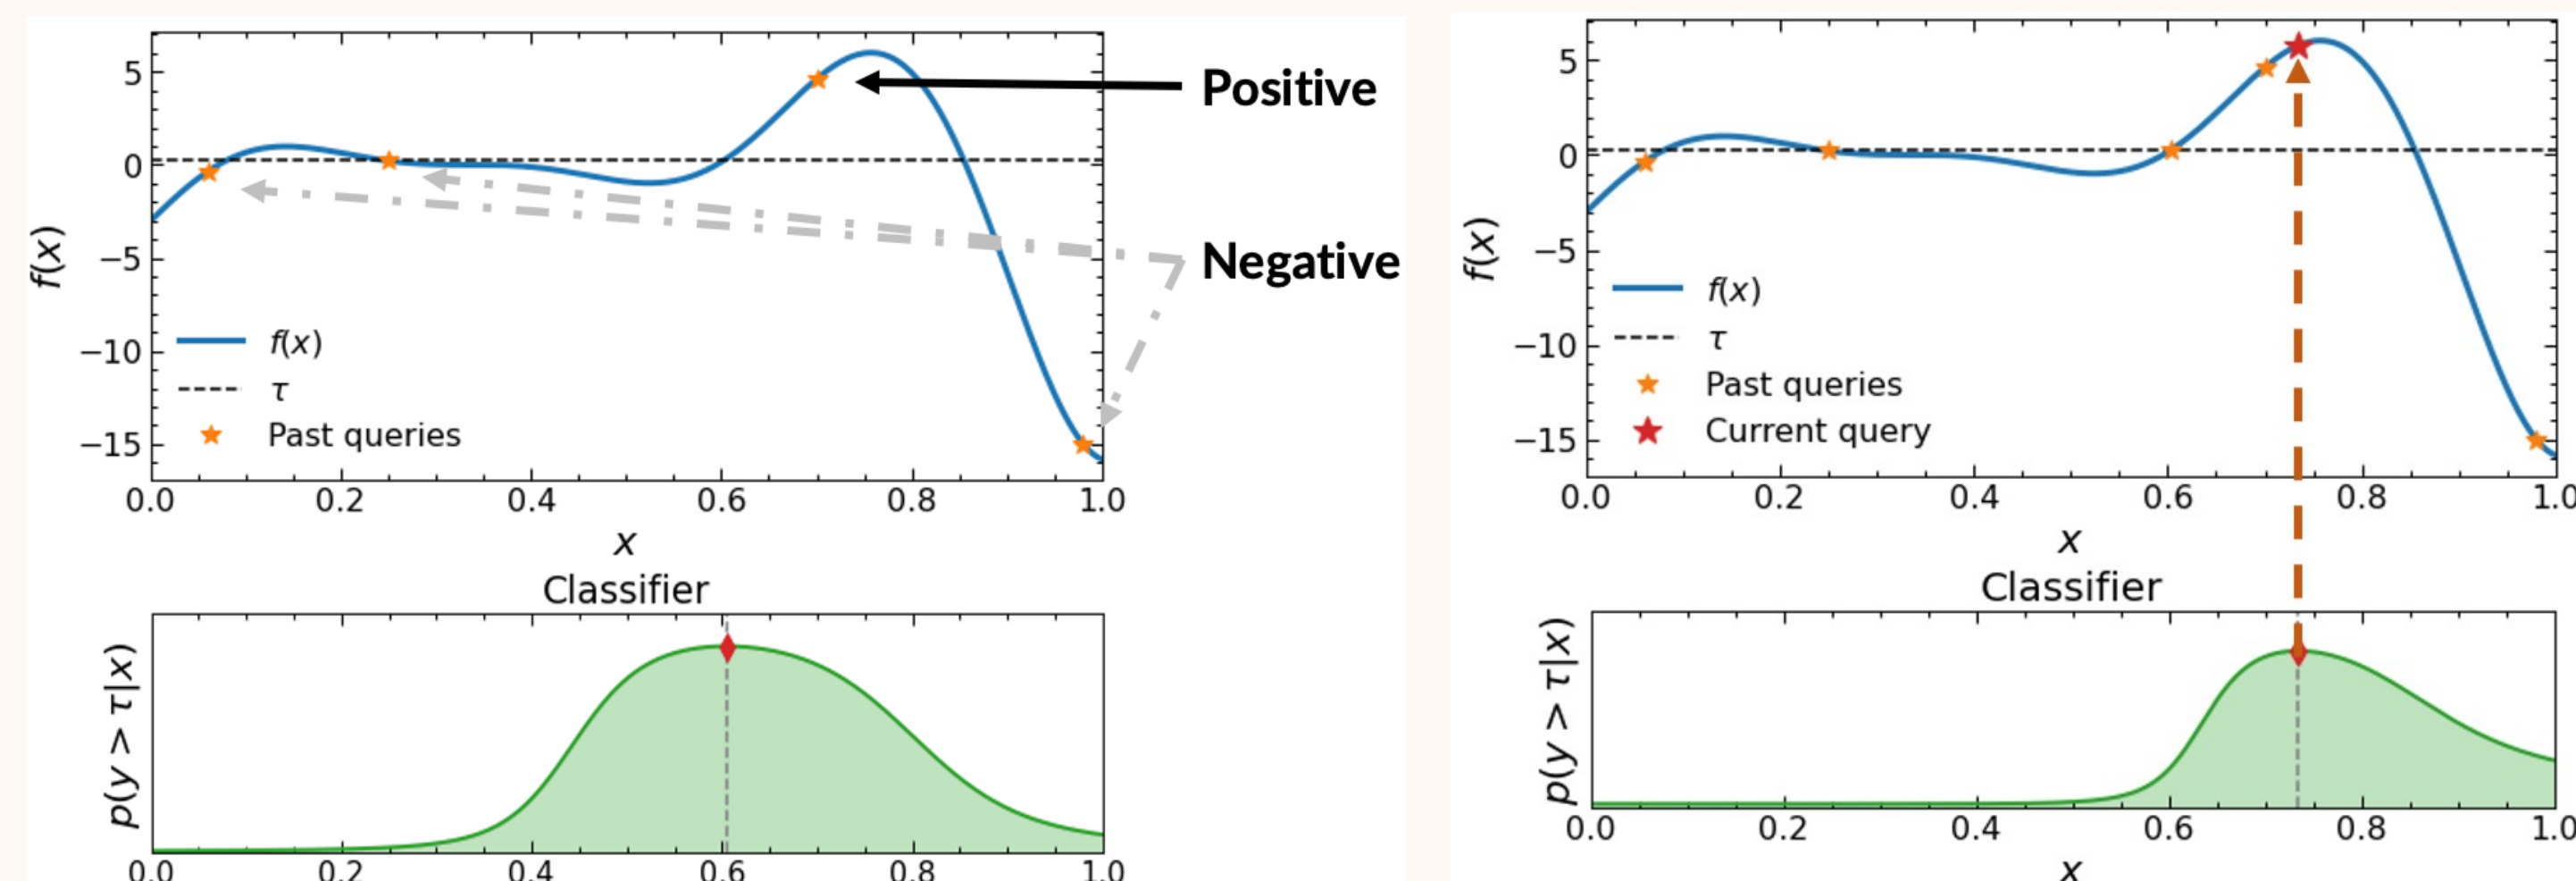
Likelihood-free Bayesian Optimization

Reweight the “positive” queries by their utility values!



Theorem: LFBO converges to desired acquisition function asymptotically.

Bayesian Optimization via Classification (BORE)



Train a classifier to predict whether a point is above the threshold.

Then use the classifier to select the next query.

This approach does not care a point is over the threshold by how much.

Only equivalent to PI

Experiments: Hyperparameter Tuning & Neural Architecture Search

